

FSDP семинар

Plan

- Prerequisites: CUDA streams / events, DeviceMesh, DTensor
- FSDP2: interface, options, internals
- PyTorch DCP, efficient garbage collection

CUDA streams and events

```
all_gather_stream = torch.cuda.Stream()

...
# layer 3 unshard
with torch.cuda.stream(all_gather_stream):
    model.layers[3].all_gather()
    all_gather_event_3 = torch.cuda.Event()
    # or all_gather_stream.record_event()

# layer 2 forward
activations = model.layers[2](activations)

# layer 4 unshard
with torch.cuda.stream(all_gather_stream):
    model.layers[4].all_gather()
    all_gather_event_4 = torch.cuda.Event()

# layer 3 forward
torch.cuda.default_stream().wait_event(all_gather_event_3)
activations = model.layers[3](activations)


...
```

CUDA streams and events



DeviceMesh

**DeviceMesh: The higher level abstraction
that manages ProcessGroups**



2 host with 2 GPUs each, represented
as a 2-D mesh [[0, 1], [2, 3]]

DeviceMesh

```
from torch.distributed.device_mesh import init_device_mesh

mesh_1d = init_device_mesh("cuda", mesh_shape=(8,), mesh_dim_names=("dp",))
mesh_2d = init_device_mesh("cpu", mesh_shape=(2, 8), mesh_dim_names=("dp", "tp"))
mesh_3d = init_device_mesh(
    "cuda",
    mesh_shape=(2, 2, 8),
    mesh_dim_names=("pp", "dp", "tp"),
)
dp_group = mesh_2d.get_group("dp")
dist.all_gather(..., group=dp_group)

mesh_2d.get_local_rank("tp")

mesh = init_device_mesh("cpu", mesh_shape=(16,), mesh_dim_names=("world",))
mesh_dp_cp = mesh._unflatten(0, mesh_sizes=(4, 4), mesh_dim_names("dp", "cp"))
mesh_host = mesh._unflatten(0, mesh_sizes=(2, 8), mesh_dim_names("inter", "intra"))
```

DTensor

```
class DTensor:
    _local_tensor: torch.Tensor
    _spec:

@dataclass
class DTensorSpec:
    mesh: DeviceMesh
    placements: tuple[Placement, ...] # (Shard(0), Replicate())

    # tensor meta will only be set during sharding propagation
    tensor_meta: TensorMeta | None = None # (dtype, shape, stride)

    def redistribute(
        placements: Sequence[Placement]
    ) -> DTensor | torch.Tensor: ...

    def __torch_dispatch__(): ...
```

DTensor

```
from torch.distributed.tensor import DTensor, distribute_tensor

mesh = init_device_mesh("cuda", mesh_shape=(8,), mesh_dim_names=("dp",))
big_tensor = torch.randn(1024, 4096)
placements = (Shard(dim=0),)

dtensor = distribute_tensor(
    big_tensor,
    device_mesh=mesh,
    placements=placements,
)
dtensor._local_tensor
dtensor.to_local() # .shape = (512, 4096)

shard = ... # .shape = (512, 4096)
dtensor = DTensor.from_local(
    shard,
    device_mesh=mesh,
    placements=placements,
) # .shape = (1024, 4096)

dtensor.redistribute(placements=(Replicate(),))
dtensor.full_tensor()
```


DTensor

- `Shard`: Tensor sharded on the tensor dimension `dim` on the devices of the `DeviceMesh` dimension
- `Replicate`: Tensor replicated on the devices of the `DeviceMesh` dimension
- `Partial`: Tensor is pending reduction on the devices of the `DeviceMesh` dimension


1. `Shard(dim) -> Replicate(): all_gather`
2. `Shard(src_dim) -> Shard(dst_dim): all_to_all`
3. `Replicate() -> Shard(dim) : local chunking (i.e. torch.chunk)`
4. `Partial() -> Replicate(): all_reduce`
5. `Partial() -> Shard(dim): reduce_scatter`

DTensor

Data: [6, 8, 10, 12]




[Shard(0)]



[Replicate()]


DeviceMesh: ([0, 1, 2, 3])



[Shard(0), Replicate()]

DeviceMesh: ([[0, 1], [2, 3]])

FSDP2



FSDP2

```
from torch.distributed.fsdp import fully_shard

mesh_2d = init_device_mesh(
    "cuda",
    mesh_shape=(2, 8),
    mesh_dim_names=("dp", "tp"),
)
model = Model()


for layer in model.layers:
    fully_shard(
        module, # (module1, module2)
        mesh=dp_mesh,
        reshard_after_forward=True, # ZeRO-3
        mp_policy=MixedPrecisionPolicy(
            param_dtype=torch.float16,
            reduce_dtype=torch.float32,
        ),
        offload_policy=CPUOffloadPolicy(),
    )

    fully_shard(model, ...)
```

```
for step in ...:
    for gas_step in ...:
        is_last_backward = gas_step == num_gas_steps - 1
        # ZeRO-2
        model.set_reshard_after_backward(is_last_backward)
        # ZeRO-1
        model.requires_gradient_sync(is_last_backward)

        loss = loss_fn(model(inputs), targets)
        ...
```

FSDP2



FSDP2 — hooks

```
register_forward_pre_hook(hook, *, prepend=False, with_kwargs=False) [SOURCE]
```

Register a forward pre-hook on the module.

The hook will be called every time before `forward()` is invoked.

```
register_forward_hook(hook, *, prepend=False, with_kwargs=False, always_call=False) [SOURCE]
```

Register a forward hook on the module.

The hook will be called every time after `forward()` has computed an output.

FSDP2 — pre-forward

```
def pre_forward(module, args):
    module.unshard() # in all-gather stream
    module.wait_for_unshard() # sync compute (default) stream with all-gather stream
    module._register_post_backward_hook(args)
    return args

def unshard(module):
    with torch.cuda.stream(all_gather_stream):
        module.all_gather_result = module.all_gather()
    module.all_gather_event = all_gather_stream.record_event()

def wait_for_unshard(module):
    torch.cuda.default_stream().wait_event(module.all_gather_event)
    module.set_unsharded_params(module.all_gather_result)

def fully_shard(module, ...):
    ...
    module.register_forward_pre_hook(pre_forward)
```

FSDP2 — post-forward

```
def post_forward(module, args, output):
    module.reshard()
    module._record_post_forward()
    module._register_pre_backward_hook(output)
    return output

def reshard(module):
    module.set_sharded_params() # and free unsharded params

def _record_post_forward(module):
    post_forward_index = len(module.comm_ctx.post_forward_order)
    module.comm_ctx.post_forward_order.append(module)
    module._post_forward_indices.append(post_forward_index)

def fully_shard(module, ...):
    ...
    module.register_forward_hook(post_forward)
```

FSDP2 — pre-backward

```
def pre_backward(module, *unused):
    module.unshard() # no-op if prefetched
    module.wait_for_unshard()
    module._backward_prefetch()

def _backward_prefetch(module):
    curr_index = module._post_forward_indices.pop()
    target_index = curr_index - 1
    target_module = self.comm_ctx.post_forward_order[target_index]
    target_module.unshard()

def _register_pre_backward_hook(self, output):
    for t in output:
        if torch.is_tensor(t) and t.requires_grad:
            t.register_hook(self._pre_backward)
    return output
```

FSDP2 — post-backward


```
def post_backward(module, *unused: Any):
    if module.reshard_after_backward:
        module.reshard()
    if module.reduce_grads:
        reduce_scatter_stream.wait_stream(torch.cuda.default_stream())
        with torch.cuda.stream(reduce_scatter_stream):
            module.reduce_scatter_grads()
        reduce_event = reduce_scatter_stream.record_event()

def _register_post_backward_hook(module, args):
    RegisterPostBackwardFunction.apply(self, *args)


class RegisterPostBackwardFunction(torch.autograd.Function):
    @staticmethod
    def forward(ctx, module, *inputs):
        ctx.module = module
        return inputs

    @staticmethod
    def backward(ctx, *grads):
        module.post_backward()
        return (None,) + grads
```

FSDP2 — memory



FSDP2 — memory




Computation / communication overlap

- Implicit prefetching
 - в `pre_forward`
- Explicit prefetching
 - в `pre_backward`
 - можно задать руками

```
module.set_modules_to_forward_prefetch(modules)
module.set_modules_to_backward_prefetch(modules)
```


Подробнее про работу со стримами




Подробнее про работу со стримами — forward




Подробнее про работу со стримами — backward



ZeRO-2




ZeRO-2



ZeRO-1



ZeRO-1



HSDP

```
mesh_2d = init_device_mesh(  
    "cpu",  
    mesh_shape=(2, 8),  
    mesh_dim_names=("dp_replicate", "dp_shard"),  
)  
  
fully_shard(  
    module,  
    mesh=mesh_2d,  
    ...  
)
```


- логика становится заметно сложнее, показывать не буду(

CPU offloading

- ZeRO-Offload

```
with torch.device("cpu"):  
    model = Model()  
  
    fully_shard(  
        module,  
        ...  
        offload_policy=CPUOffloadPolicy(),  
    )  
  
def unshard(module):  
    sharded_param = sharded_param.to(  
        device,  
        non_blocking=True,  
    )  
    ...  
    module.all_gather()  
  
def post_backward(module):  
    new_sharded_grad = new_sharded_grad.to(  
        torch.device("cpu"),  
        non_blocking=True  
    )
```

CPU offloading




hpZ


- ZeRO++

```
mesh = init_device_mesh(  
    "cuda",  
    mesh_shape=(16,),  
    mesh_dim_names=("dp",),  
)  
fully_shard(  
    module,  
    mesh,  
    ...  
    reshard_after_forward=8,  
)
```

hpZ



hpZ



PyTorch DCP

- два вида `state_dict`
 - `SHARDED_STATE_DICT`
 - `FULL_STATE_DICT`
- в FSDP2 всегда sharded, но состоит из DTensor-ов
 - с помощью `.redistribute()` можно менять шардирование чекпоинта
- DCP умеет эффективно отгружать чекпоинты с минимальным оверхедом


PyTorch DCP

```
import torch.distributed.checkpoint as dcp
model = Model()
fully_shard(model)
optimizer = Optimizer(model.parameters())


state_dict = {
    "model": model.state_dict(),
    "optimizer": optimizer.state_dict()
}
dcp.state_dict_saver.save(state_dict)
dcp.state_dict_loader.load(state_dict)
```

- truthfully it's a bit more complicated


PyTorch DCP




PyTorch DCP



PyTorch DCP



PyTorch DCP



Garbage collection tuning

```
gc.disable()  
gc.collect(1)  
  
... init  
  
for step in ...:  
    if step > 1 and step % _gc_freq == 0:  
        gc.collect(1)  
  
    ... step
```

Extras

- SimpleFSDP
- `unshard_in_backward`
- meta device init
- compile